



東南大學  
SOUTHEAST UNIVERSITY

# OPERATING SYSTEM CONCEPTS

.....

## Chapter 11. File-System Interface

A/Prof. Kai Dong

# Objectives



- To explain the function of file systems.
- To describe the interfaces to file systems.
- To discuss file-system design tradeoffs, including access methods, file sharing, file locking, and directory structures.
- To explore file-system protection.



# Contents

1. Files and Directories
2. File Interface
3. Directory Interface
4. Links
5. File System Interface



# Contents

1. Files and Directories
2. File Interface
3. Directory Interface
4. Links
5. File System Interface



# Files and Directories

- Virtualizing the persistent storage
- Two key abstractions
  - A **file** is simply a linear array of bytes, each of which you can read or write
    - » Each file has some kind of low-level name, which is often referred to as its **inode number**
  - A **directory**, contains a list of entries, each of which is a (user-readable name, low-level name) pair, referring to either a file or other directory
    - » Also has a low-level name (i.e., an inode number)



## Files and Directories

- By placing directories within other directories, users are able to build an arbitrary **directory tree** (or directory hierarchy), under which all files and directories are stored.
- The directory hierarchy starts at a **root directory**, and uses some kind of **separator** to name subsequent **sub-directories** until the desired file or directory is named.
- A file can thus be referred to by its absolute pathname.
- One great thing provided by the file system: a convenient way to **name** all the files we are interested in.
- The file name often has two parts: x.y, separated by a period. The first part is an arbitrary name, whereas the second part of the file name is usually used to indicate the **type** of the file.
  - However, this is usually just a convention: there is usually no enforcement that the data contained in a file named main.c is indeed C source code.



# Contents

1. Files and Directories
2. File Interface
3. Directory Interface
4. Links
5. File System Interface



# File Interface

## Creating Files

- The `open()` system call

```
1  int fd = open("foo", O_CREAT | O_WRONLY | O_TRUNC);
```

- Creates a file called "foo" in the current working directory.
- The routine `open()` takes a number of different flags.
  - The program creates the file (`O_CREAT`), can only write to that file while opened in this manner (`O_WRONLY`), and, if the file already exists, first truncate it to a size of zero bytes thus removing any existing content (`O_TRUNC`).





## File Interface

### Creating Files (contd.)

- A **file descriptor** is what `open()` returns.
- Is just an integer, private per process, and is used in UNIX systems to access files, i.e., to call `read()` and `write()`.
- We will see how to use a file descriptor.



# File Interface

## Reading Files

- Reading an existing file

```
1 prompt> echo hello > foo
2 prompt> cat foo
3 hello
4 prompt>
```

- Let's trace the system calls

```
1 prompt> strace cat foo
2 ...
3 open ("foo", O_RDONLY | O_LARGEFILE) = 3
4 read(3, "hello\n", 4096) = 6
5 write(1, "hello\n", 6) = 6
6 hello
7 read(3, "", 4096) = 0
8 close(3) = 0
9 ...
10 prompt>
```



## File Interface

### Reading Files (contd.)

```
1 open ("foo", O_RDONLY | O_LARGEFILE) = 3
```

- Each running process already has three files open, standard input, standard output, and standard error.

```
1 read(3, "hello\n", 4096) = 6
```

- The first argument to *read()* is the file descriptor, thus telling the file system which file to read;
- The second argument points to a buffer where the result of the *read()* will be placed (the system-call trace above shows the results of the read in this spot).
- The third argument is the size of the buffer, which in this case is 4 KB.
- The call to *read()* returns the number of bytes it read ( take into account the end-of-line marker).



## File Interface

### Reading Files (contd.)

```
1 write(1, "hello\n", 6) = 6
```

- The file descriptor 1 is the standard output, and thus is used to write the word “hello” to the screen as the program cat is meant to do.

```
1 read(3, "", 4096) = 0
```

- The cat program then tries to read more from the file, but since there are no bytes left in the file, the `read()` returns 0 and the program knows that this means it has read the entire file.

```
1 close(3) = 0
```

- Close the file.

# File Interface

## Writing Files



- Writing a file has similar steps
  - First, a file is opened for writing;
  - Then the `write()` system call is called;
  - Perhaps repeatedly for larger files;
  - And then `close()`.



## File Interface

### Reading / Writing NOT Sequentially

- Thus far, we've discussed how to read and write files, but all access has been sequential.
- How to read or write to a specific offset within a file.
- The *lseek()* system call

```
1 off_t lseek(int fd, off_t offset, int whence);
```

- The first argument is a file descriptor.
- The second argument is *offset*, which positions the file offset to a particular location within the file.
- The third argument, called *whence* for historical reasons, determines exactly how the seek is performed.



## File Interface

### Reading / Writing NOT Sequentially (contd.)

- Details about *whence*

- 1 If *whence* is `SEEK_SET`, the offset is set to offset bytes.
- 2 If *whence* is `SEEK_CUR`, the offset is set to its current location plus offset bytes.
- 3 If *whence* is `SEEK_END`, the offset is set to the size of the file plus offset bytes.

- Thus, we know that an open file has a current offset, which is updated in one of two ways.
  - The first is when a read or write of *N* bytes takes place, *N* is added to the current offset; thus each read or write **implicitly** updates the offset.
  - The second is **explicitly** with *lseek*, which changes the offset as specified above.

# File Interface

## Writing Immediately



- Most times when a program calls `write()`, it is just telling the file system: please write this data to persistent storage, at some point in the future.
- The file system, for performance reasons, will **buffer** such writes in memory for some time, then the write(s) will actually be issued to the storage device.
- How to write immediately?





## File Interface

### Writing Immediately (contd.)

- The `fsync()` system call.
- When a process calls `fsync(int fd)`, the file system responds by forcing all **dirty** (i.e., not yet written) data to disk, for the file referred to by the specified file descriptor.

```
1  int fd = open("foo", O_CREAT | O_WRONLY | O_TRUNC);  
2  assert(fd > -1);  
3  int rc = write(fd, buffer, size);  
4  assert(rc == size);  
5  rc = fsync(fd);  
6  assert (rc == 0);
```

- If the file is newly created, we also need to `fsync()` the directory that contains the file.
- What if something bad (power off) happens when performing `fsync()`?



# File Interface

## *Renaming Files*

- How to give a file a different name.

```
1 prompt> mv foo bar
```

- The *rename()* system call

```
1 rename(char *old, char *new)
```

- the original name of the file (*old*)
  - the new name (*new*).
- It is implemented as an **atomic** call with respect to system crashes.
- It is critical for supporting certain kinds of applications that require an atomic update to file state.



# File Interface

## Renaming Files (contd.)

- Atomic update to file state

```
1  int fd = open("foo.txt.tmp", O_WRONLY | O_CREAT | O_TRUNC);
2  write(fd, buffer, size);
3  fsync(fd);
4  close(fd);
5  rename("foo.txt.tmp", "foo.txt");
```



## File Interface

### Getting Information about Files

- File system keeps a fair amount of information about each file it is storing.
- Such data is called **metadata**.

```
1  struct stat {  
2      dev_t  st_dev;           // ID of device containing file  
3      ino_t  st_ino;          // inode number  
4      mode_t st_mode;         // protection  
5      nlink_t      st_nlink;    // number of hard links  
6      uid_t  st_uid;          // user ID of owner  
7      gid_t  st_gid;          // group ID of owner  
8      dev_t  st_rdev;         // device ID (if special file)  
9      off_t  st_size;         // total size, in bytes  
10     blksize_t      st_blksize; // blocksize for filesystem I/O  
11     blkcnt_t      st_blocks;   // number of blocks allocated  
12     time_t        st_atime;    // time of last access  
13     time_t        st_mtime;    // time of last modification  
14     time_t        st_ctime;    // time of last status change  
15 }
```



## File Interface

### Getting Information about Files (contd.)

- The `stat()` or `fstat()` system calls

```
1 prompt> echo hello > file
2 prompt> stat file
3   File: 'file'
4   Size: 6 Blocks: 8      IO Block: 4096   regular file
5   Device: 811h/2065d    Inode: 67158084  Links:1
6   Access: (0640/-rw-r-----)  Uid: (30686/  kai)  Gid: (30686/  kai)
7   Access: (2019-07-13) 15:50:20.157594748 -0500
8   Modify: (2019-07-13) 15:50:20.157594748 -0500
9   Change: (2019-07-13) 15:50:20.157594748 -0500
```

- Such information is kept in a structure called an **inode**, we will dive into it in future.



# File Interface

## Permission Bits

- UNIX **permission bits**

```
1 prompt> echo hello > file
2 prompt> ls -l foo.txt
3 -rw-r--r-- 1 kai kaigroup 0 Jul 13 16:29 foo.txt
```

- Nine bits determine, for each regular file, directory, and other entities, exactly who (**owner**, **group**, **other**) can access it and how (**read**, **write**, **execute**).

- Change the file mode (*chmod()* command)

```
1 prompt> chmod 777 test.sh
```



# File Interface

## Removing Files

- Delete files: *rm()* in UNIX.
- Let's trace the system calls in program *rm*.

```
1  prompt> strace rm foo
2  ...
3  unlink("foo")                = 0
4  ...
```

- The *unlink()* system call
  - Takes the name of the file to be removed.
  - Why not something like *remove()* or *delete()*?
    - » To understand the answer to this puzzle, we must first understand more than just files, but also directories.



# Contents

1. Files and Directories
2. File Interface
- 3. Directory Interface**
4. Links
5. File System Interface





# Directory Interface

## Making Directories

- The *mkdir()* system call

```
1 prompt> strace mkdir foo
2 ...
3 mkdir("foo", 0777)                = 0
4 ...
5 prompt>
```

- A newly created directory is considered “empty”, (i.e., only has “.” and “..” entries).

```
1 prompt> ls -a
2 ./      ../
3 prompt> ls -al
4 total 8
5 drwxr-x--- 2 kai kai 6 Jul 13 16:17 ./
6 drwxr-x--- 26 kai kai 4096 Jul 13 16:17 ../
7 prompt>
```



# Directory Interface

## Reading Directories

- The `ls()` system call

```
1  /* a ls() like program */
2  struct dirent {
3      char          d_name[256];    // filename
4      ino_t         d_no;           // inode number
5      off_t         d_off;          // offset to the next dirent
6      unsigned short d_reclen;      // length of this record
7      unsigned char  d_type;        // type of file
8  };
9
10 int main(int argc, char *argv[]) {
11     DIR *dp = opendir(".");
12     assert(dp != NULL);
13     struct dirent *d;
14     while ((d = readdir(dp)) != NULL) {
15         printf("%lu %s\n", (unsigned long) d->d_ino, d->d_name
16             );
17     }
18     closedir(dp);
19     return 0;
20 }
```

# Directory Interface

## Deleting Directories



- The program *rmdir*
- The *rmdir()* system call
- *rmdir()* has the requirement that the directory be empty (i.e., only has “.” and “..” entries) before it is deleted.



# Contents

1. Files and Directories
2. File Interface
3. Directory Interface
4. **Links**
5. File System Interface



# Links

## Hard Links

- The `link()` system call takes two arguments, an old pathname and a new one; when you “link” a new file name to an old one, you essentially create another way to refer to the same file.

```
1  prompt> echo hello > file
2  prompt> cat file
3  hello
4  prompt> ln file file2
5  prompt> cat file2
6  hello
7  prompt> ls -i file file2
8  67158084 file
9  67158084 file2
10 prompt>
```

- The way link works is that it simply creates another name in the directory you are creating the link to, and refers it to the same inode number (i.e., low-level name) of the original file. The file is **NOT** copied in any way.



## Links

### Hard Links (contd.)

- Recall that removing a file is performed via *unlink()*.
- On creating a file
  - First, you are making a structure (the inode) that will track virtually all relevant information about the file, including its size, where its blocks are on disk, and so forth.
  - Second, you are linking a human-readable name to that file, and putting that link into a directory.
- When unlinking a file
  - The file system checks a reference count (sometimes called the link count) within the inode number.
  - Only when the reference count reaches zero, does the file system also free the inode and related data blocks, and thus truly “delete” the file.



# Links

## Symbolic Links

- **Hard links** are somewhat limited.
  - You can't create one to a directory (for fear that you will create a cycle in the directory tree);
  - You can't hard link to files in other disk partitions (because inode numbers are only unique within a particular file system, not across file systems);
  - Etc.
- There is one other type of link that is really useful, and it is called a **symbolic link** or sometimes a soft link.

```
1 prompt> echo hello > file
2 prompt> ln -s file file2
3 prompt> cat file2
4 hello
5 prompt>
```



# Links

## Symbolic Links

- A symbolic link is actually a file itself, of a different type.

```
1 prompt> stat file
2 ... regular file ...
3 prompt> stat file2
4 ... symbolic link ...
```

- Running `ls`

```
1 prompt> ls -al
2 drwxr-x--- 2 kai kai 6 Jul 13 19:10 ./
3 drwxr-x--- 26 kai kai 4096 Jul 13 16:17 ../
4 -rw-r----- 1 kai kai 6 Jul 13 19:10 file
5 lrwxrwxrwx 1 kai kai 4 Jul 13 19:10 file2 -> file
```

- The first character in the left-most column is a “-” for regular files, a “d” for directories, and an “l” for soft links.
- The size of the symbolic link is 4 bytes in this case.





# Links

## Symbolic Links

- A dangling reference:

```
1 prompt> echo hello > file
2 prompt> ln -s file file2
3 prompt> cat file2
4 hello
5 prompt> rm file
6 prompt> cat file2
7 cat: file2: No such file or directory
```

- Removing the original file causes the link to point to a pathname that no longer exists.



# Contents

1. Files and Directories
2. File Interface
3. Directory Interface
4. Links
5. File System Interface



# File System Interface

## Making and Mounting a File System

- How to assemble a full directory tree from many underlying file systems?
  - Make file systems (*mkfs()* command)
    - » Give the tool, as input, a device (such as a disk partition, e.g., `/dev/sda1`) and a file system type (e.g., `ext3`), and it simply writes an empty file system, starting with a root directory, onto that disk partition.
  - Mount them to make their contents accessible (*mount()* command)
    - » Takes an existing directory as a target mount point and essentially paste a new file system onto the directory tree at that point.

```
1 prompt> mount -t ext3 /dev/sda1 /home/users
```